



Latent-space robustness: embedding sensitivity vs descriptor-space separation (Cr-doped Sb_2Te_3 , 600 K)


(a) t-SNE $p=5$, seed=0 — sil=0.36, trust=0.97




(b) t-SNE $p=5$, seed=1 — sil=0.36, trust=0.97




(c) t-SNE $p=5$, seed=2 — sil=0.36, trust=0.97




(d) t-SNE $p=30$, seed=0 — sil=0.41, trust=0.97




(e) t-SNE $p=30$, seed=1 — sil=0.41, trust=0.97




(f) t-SNE $p=30$, seed=2 — sil=0.41, trust=0.97




(g) t-SNE $p=50$, seed=0 — sil=0.42, trust=0.96



(h) t-SNE $p=50$, seed=1 — sil=0.42, trust=0.96



(i) t-SNE $p=50$, seed=2 — sil=0.42, trust=0.96



(j) Overall silhouette by space (error bars = std across perplexity×seed settings)

